

MACHINE LEARNING

Liviu Ciortuz

Department of CS, University of Iași, România

What is Machine Learning?

- ML studies algorithms that improve with experience.
learn from

Tom Mitchell's Definition of the [general] learning problem:

“A computer program is said to *learn* from experience E with respect to some class of *tasks* T and *performance measure* P , if its performance on tasks in T , as measured by P , improves with experience E .”

- Examples of [specific] learning problems (see next slide)
- [Liviu Ciortuz:] ML is data-driven programming
- [Liviu Ciortuz:] ML gathers a number of well-defined sub-domains/**disciplines**, each one of them aiming to solve in its own way the above-formulated [general] learning problem.

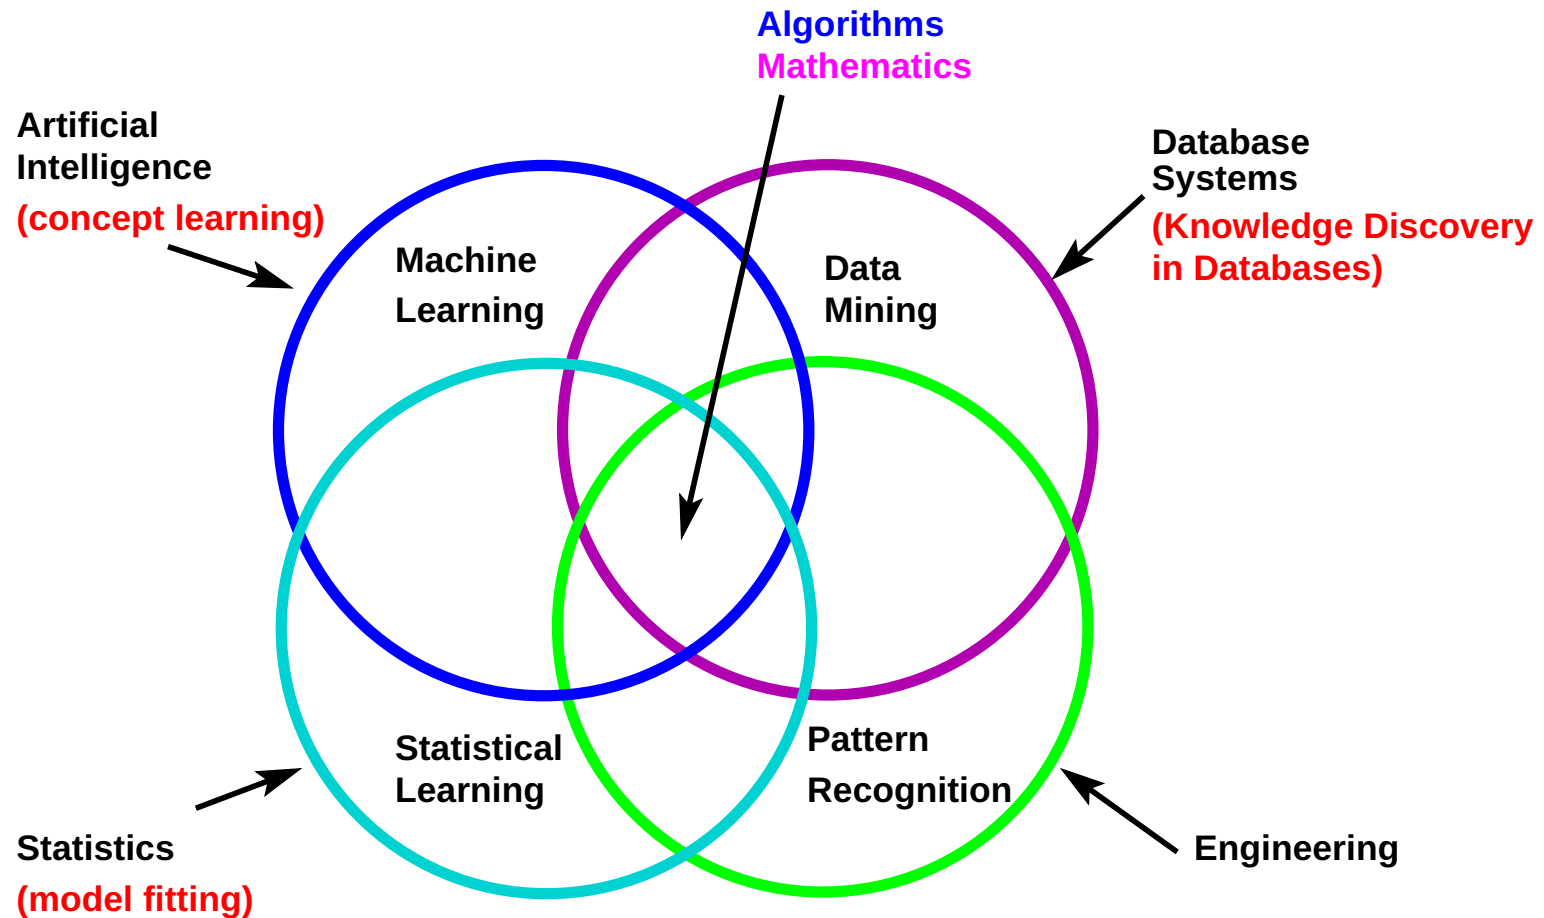
What is Machine Learning good for?

- natural language (text & speech) processing
- genetic sequence analysis
- robotics
- customer (financial risk) evaluation
- terrorist threat detection
- compiler optimisation
- semantic web
- computer security
- software engineering
- computer vision (image processing)
- etc.

Related courses at FII

- Artificial Intelligence
 - Genetic Algorithms
 - Artificial Neural Networks
 - Probabilistic programming
-
- Special Chapters of Machine Learning
 - Special Chapters of Artificial Intelligence
 - Special Chapters of Artificial Neural Networks
 - Data Mining
 - Nature-inspired computing methods
 - Big Data Analytics
 - Image Processing
 - Computer Vision
-
- Bioinformatics

A multi-domain view



The Machine Learning Undergraduate Course: Plan

0. Introduction to Machine Learning (T. Mitchell, ch. 1)

1. **Probabilities Revision** (Ch. Manning & H. Schütze, ch. 2)

2. Decision Trees (T. Mitchell, ch. 3)

3. Bayesian Learning (T. Mitchell, ch. 6)
[and the relationship with Logistic Regression]

4. Instance-based Learning (T. Mitchell, ch. 8)

5. Clustering Algorithms (Ch. Manning & H. Schütze, ch. 14)

The Machine Learning Master Course:

Tentative Plan

1. Decision Trees: Boosting
2. Support Vector Machines (N. Cristianini & J. Shawe-Taylor, 2000)
3. Computational Learning Theory (T. Mitchell, ch. 7)

Probabilities Revision (Ch. Manning & H. Schütze, ch. 2)

4. Gaussian Bayesian Learning
5. The EM algorithmic schemata (T. Mitchell, ch. 6.12)
6. Hidden Markov Models (Ch. Manning & H. Schütze, ch. 9)

Bibliography

0. “Exerciții de învățare automată”

L. Ciortuz, A. Munteanu E. Bădăraș.

Iași, Romania, 2023

www.info.uaic.ro/~ciortuz/ML.ex-book/editia-2023f/ex-book.20sept2023.pdf

1. “Machine Learning”

Tom Mitchell. McGraw-Hill, 1997

2. “Machine Learning Foundations”

Teaho Jo. Springer, 2021

3. “Deep Machine Learning Foundations”

Teaho Jo. Springer, 2023

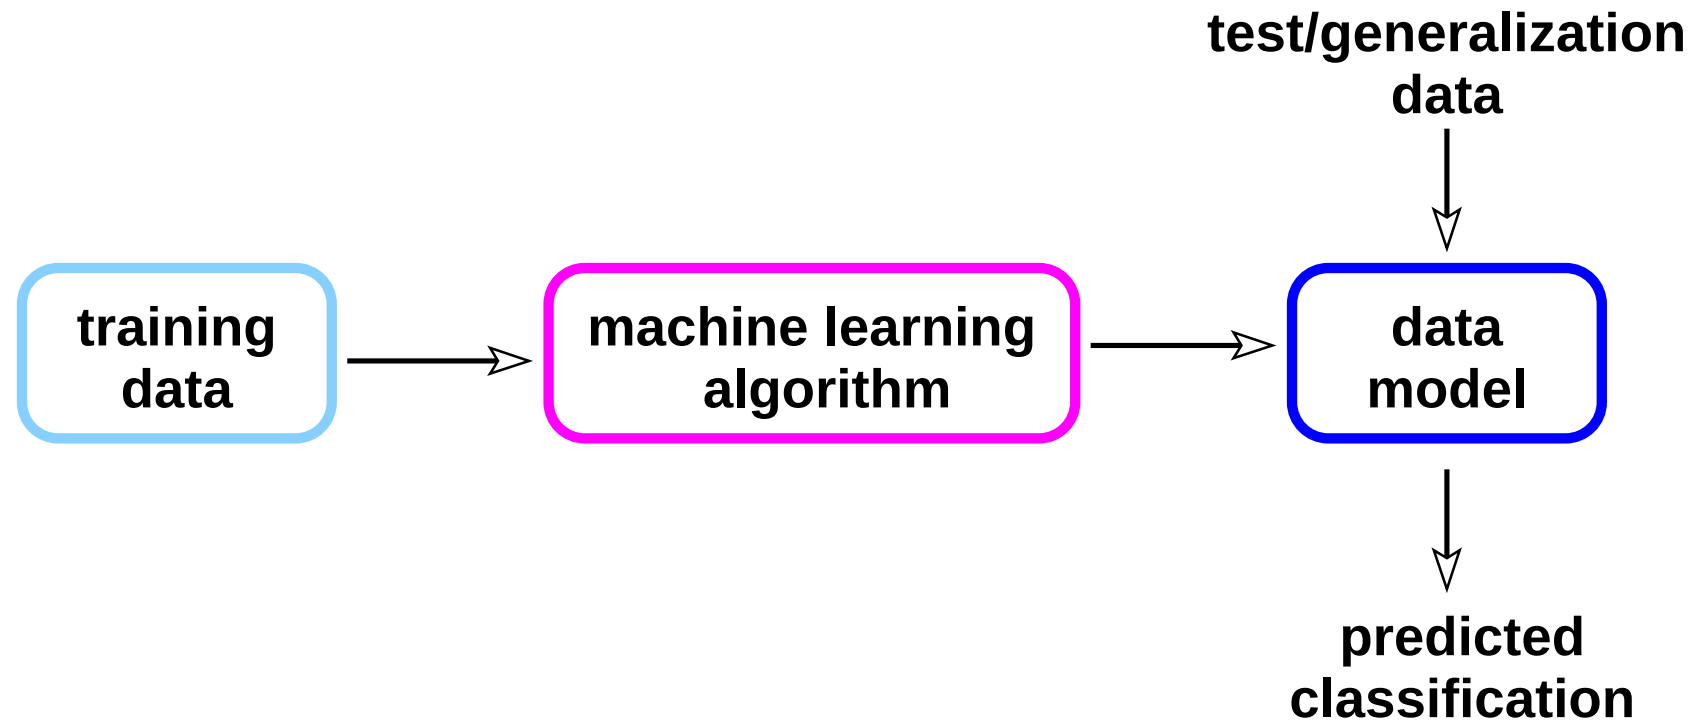
4. “Foundations of Statistical Natural Language Processing”

Christopher Manning, Hinrich Schütze. MIT Press, 2002

5. “Support Vector Machines and other kernel-based learning methods”

Nello Cristianini, John Shawe-Taylor. Cambridge University Press, 2000.

A general schema for machine learning methods



*“We are drawing in **information** but starved for **knowledge**.”*

John Naisbitt, “Megatrends” book, 1982

Basic ML Terminology

1. instance x , instance set X
concept $c \subseteq X$, or $c : X \rightarrow \{0, 1\}$
example (labeled instance): $\langle x, c(x) \rangle$; positive examples, neg. examples
2. hypotheses $h : X \rightarrow \{0, 1\}$
hypotheses representation language
hypotheses set H
hypotheses consistent with the concept c : $h(x) = c(x), \forall$ example $\langle x, c(x) \rangle$
version space
3. learning = train + test
supervised learning (classification), unsupervised learning (clustering)
4. $error_h = | \{x \in X, h(x) \neq c(x)\} |$
training error, test error
accuracy, precision, recall
5. validation set, development set
 n -fold cross-validation, leave-one-out cross-validation
overfitting

The Inductive Learning Assumption

Any hypothesis found to conveniently approximate the target function over a sufficiently large set of training examples

will also conveniently approximate the target function over other unobserved examples.

Inductive Bias

Consider

- a concept learning algorithm L
- the instances X , and the target concept c
- the training examples $D_c = \{\langle x, c(x) \rangle\}$.
- Let $L(x_i, D_c)$ denote the classification assigned to the instance x_i by L after training on data D_c .

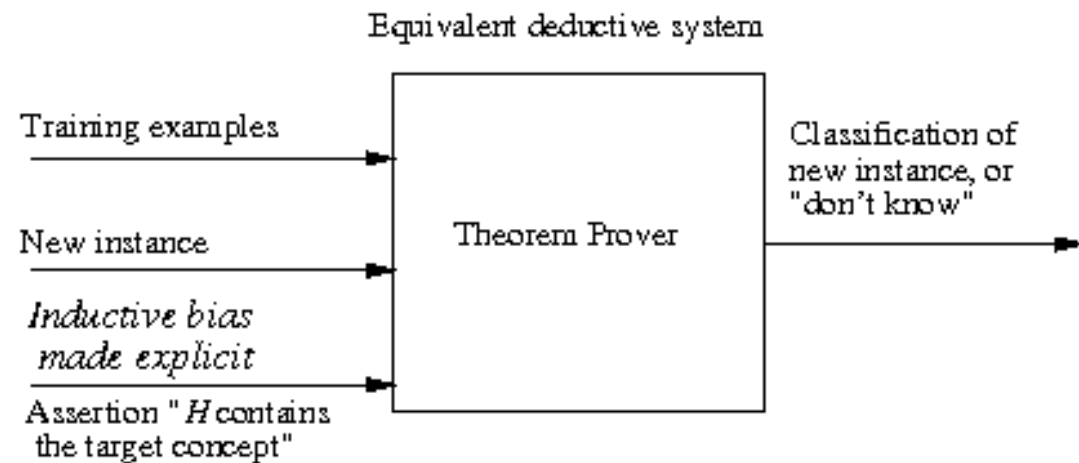
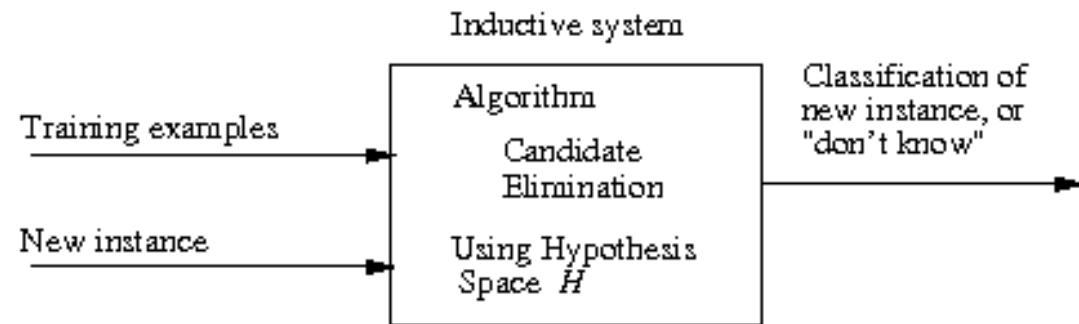
Definition:

The **inductive bias** of L is any minimal set of assertions B such that

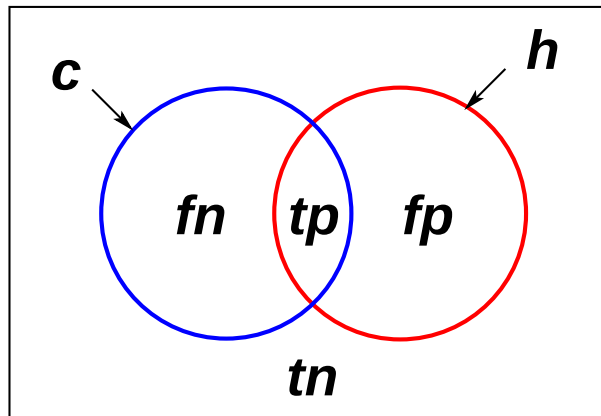
$$(\forall x_i \in X)[(B \vee D_c \vee x_i) \vdash L(x_i, D_c)]$$

for any target concept c and corresponding training examples D_c .
 ($A \vdash B$ means A logically entails B)

Inductive systems
can be modelled by
equivalent deductive
systems



Evaluation measures in Machine Learning



tp – true positives
 fp – false positives
 tn – true negatives
 fn – false negatives

accuracy: $Acc = \frac{tp + tn}{tp + tn + fp + fn}$

precision: $P = \frac{tp}{tp + fp}$

recall (or: sensitivity): $R = \frac{tp}{tp + fn}$

F-measure: $F = \frac{2 P \times R}{P + R}$

specificity: $Sp = \frac{tn}{tn + fp}$

follout: $= \frac{fp}{tn + fp}$

Mathew's Correlation Coefficient:

$$MCC = \frac{tp \times tn - fp \times fn}{\sqrt{(tp + fp) \times (tn + fn) \times (tp + fn) \times (tn + fp)}}$$

Lazy learning vs. eager learning algorithms

Eager: generalize before seeing query

- ID3, Backpropagation, Naive Bayes, Radial basis function networks, ...
- Must create global approximation

Lazy: wait for query before generalizing

- k -Nearest Neighbor, Locally weighted regression, Case based reasoning
- Can create many local approximations

Does it matter?

If they use the same hypothesis space H , lazy learners can represent **more complex functions**.

E.g., a lazy Backpropagation algorithm can learn a NN which is different for each query point, compared to the eager version of Backpropagation.

Basic Machine Learning Algorithms

ID3 algorithm: a simplified version

Ross Quinlan, 1979, 1986

START

create the root *node*;
assign all examples to the root node;

Main loop:

1. $A \leftarrow$ the “best” decision attribute for the next *node*;
2. for each value of A , create a new descendant of *node*;
3. sort training examples to leaf nodes;
4. if training examples are perfectly classified, then
STOP;
else iterate over the new leaf nodes

AdaBoost algorithm [Yoav Freund, Robert Schapire, 1996, 1997, 1999]

Consider m training examples $S = \{(x_1, y_1), \dots, (x_m, y_m)\}$, where $x_i \in \mathcal{X}$ and $y_i \in \{-1, +1\}$.

Suppose we have a *weak learning algorithm* A which produces a hypothesis $h : \mathcal{X} \rightarrow \{-1, +1\}$ given any distribution D of examples.

- Begin with a uniform **distribution** $D_1(i) = \frac{1}{m}$, $i = 1, \dots, m$.
- At each **iteration** $t = 1, \dots, T$,
 - run the weak learning algo A on the distribution D_t and produce the **hypothesis** h_t ;

Note (1): Since A is a weak learning algorithm, the produced hypothesis h_t at round t is only slightly better than random guessing, say, by a margin γ_t :

$$\varepsilon_t = \text{err}_{D_t}(h_t) = \Pr_{x \sim D_t}[y \neq h_t(x)] = \frac{1}{2} - \gamma_t.$$

Note (2): If at a certain iteration $t \leq T$ the weak classifier A cannot produce a hypothesis better than random guessing (i.e., $\gamma_t = 0$) or it produces a hypothesis for which $\varepsilon_t = 0$, then the AdaBoost algorithm should be stopped.

- update the **distribution**

$$D_{t+1}(i) = \frac{1}{Z_t} \cdot D_t(i) \cdot e^{-\alpha_t y_i h_t(x_i)} \quad \text{for } i = 1, \dots, m, \quad (1)$$

where $\alpha_t \stackrel{\text{not.}}{=} \frac{1}{2} \ln \frac{1 - \varepsilon_t}{\varepsilon_t}$, and Z_t is the **normalizer**.

- In the end, **deliver** $H_T = \text{sign}\left(\sum_{t=1}^T \alpha_t h_t\right)$ as the learned hypothesis, which will act as a *weighted majority vote*.

AdaBoost as an instance of a more general *stepwise algorithm*

Input: S, T, \mathcal{H}, ϕ , where

$S = \{(x_1, y_1), \dots, (x_m, y_m)\}$ is the training dataset, with

$y_i \in \{-1, +1\}$

T is the number of iterations to be executed,

\mathcal{H} is a set of “hypotheses”,

$\phi(y, y')$ is a “loss” / “cost” / “risk” function;

Procedure:

Initialize the classifier by taking $f_0(x) = 0$ (the constant function 0),

and $D_1(i) = 1/m$ for $i = 1, \dots, m$

for $t = 1$ to T do:

1. Compute

$$(h_t, \alpha_t) = \arg \min_{\alpha \in \mathbb{R}, h \in \mathcal{H}} \sum_{i=1}^m \phi(y_i, f_{t-1}(x_i) + \alpha h(x_i))$$

2. Update the classifier

$$f_t(x) = f_{t-1}(x) + \alpha_t h_t(x)$$

compute D_{t+1}

end for

return the classifier $\text{sign}(f_T(x))$

Note: At each step, the algorithm greedily adds a hypothesis $h \in \mathcal{H}$ to the current *combined hypothesis* to minimize the ϕ -loss.

A Generalized AdaBoost Algorithm

[MIT, 2003 fall, Tommi Jaakkola, HW4, pr. 2.1-3]

Initialization: $\tilde{W}_i^{(1)} = 1/m$ and $f_0(x_i) = 0$ for $i = 1, \dots, m$.

Loop: for $t = 1$ to T do:

Step 1: Find a classifier $h(x; \hat{\theta}_t)$ performing better than chance wrt the weighted training error:

$$\varepsilon_t \stackrel{\text{not.}}{=} \sum_{i: y_i \neq h(x_i; \hat{\theta}_t)} \tilde{W}_i^{(t)} y_i h(x_i; \theta) = \frac{1}{2} \left(1 - \sum_{i=1}^m \tilde{W}_i^{(t)} y_i h(x_i; \hat{\theta}_t) \right).$$

Note: Minimizing ε_t is equivalent to finding $\hat{\theta}_t$ that minimizes $\frac{\partial}{\partial \alpha} J_t(\alpha, \theta_t)|_{\alpha=0}$ wrt θ_t , where

$$J_t(\alpha, \theta_t) = \frac{1}{m} \sum_{i=1}^m \text{Loss}(y_i f_{t-1}(x_i) + y_i \alpha h(x_i; \theta_t)).$$

Step 2: Set the votes α_t for the new component by minimizing the overall empirical loss:

$$\alpha_t = \arg \min_{\alpha \geq 0} J_t(\alpha, \hat{\theta}_t).$$

Step 3: Recompute the normalized weights for the next iteration according to

$$\tilde{W}_i^{(t+1)} = -c_t \cdot \underbrace{dL(y_i f_{t-1}(x_i) + y_i \alpha_t h(x_i; \hat{\theta}_t))}_{y_i f_t(x_i)} \quad \text{for } i = 1, \dots, m,$$

where c_t is chosen so that $\sum_{i=1}^m \tilde{W}_i^{(t+1)} = 1$.

Output: f_T

The Naive Bayes Classifier

- Assume that the attributes $\langle a_1, \dots, a_n \rangle$ that describe instances are conditionally independent w.r.t. to the given classification:

$$P(a_1, a_2 \dots a_n | v_j) = \prod_i P(a_i | v_j)$$

- Training procedure:

NAIVE_BAYES_LEARN(*examples*)

for each value v_j of the output attribute

$\hat{P}(v_j) \leftarrow$ estimate $P(v_j)$

for each value a_i of each input attribute a

$\hat{P}(a_i | v_j) \leftarrow$ estimate $P(a_i | v_j)$

- The *decision rule* of the Naive Bayes classifier is:

$$\begin{aligned} v_{MAP} &= \operatorname{argmax}_{v_j \in V} P(v_j | a_1, a_2 \dots a_n) = \operatorname{argmax}_{v_j \in V} \frac{P(a_1, a_2 \dots a_n | v_j) P(v_j)}{P(a_1, a_2 \dots a_n)} \\ &= \operatorname{argmax}_{v_j \in V} P(a_1, a_2 \dots a_n | v_j) P(v_j) = \operatorname{argmax}_{v_j \in V} \prod_i P(a_i | v_j) P(v_j) \stackrel{\text{not.}}{=} v_{NB} \end{aligned}$$

Logistic Regression

Given the dataset $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(n)}, y^{(n)})\}$, where each vector $x^{(i)}$ has d features / attributes, and $y^{(i)} \in \{0, 1\}$ for $i = 1, \dots, n$, its complete *log-likelihood* is:

$$\begin{aligned} \text{log-likelihood} &= \ln \prod_{i=1}^n P(x^{(i)}, y^{(i)}) = \ln \prod_{i=1}^n (P_{Y|X}(y^{(i)}|x^{(i)}) P_X(x^{(i)})) \\ &= \ln \left(\left(\prod_{i=1}^n P_{Y|X}(y^{(i)}|x^{(i)}) \right) \cdot \left(\prod_{i=1}^n P_X(x^{(i)}) \right) \right) \\ &= \ln \prod_{i=1}^n P_{Y|X}(y^{(i)}|x^{(i)}) + \ln \prod_{i=1}^n P_X(x^{(i)}) \stackrel{\text{not.}}{=} \ell(w) + \ell_x. \end{aligned}$$

Note that ℓ_x does not depend on the parameter w .

It can be shown that the conditional *log-likelihood* function $\ell(w)$ can be written as:

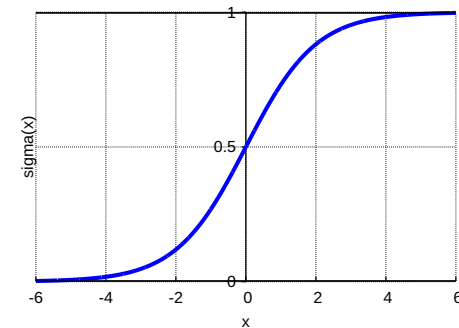
$$\begin{aligned} \ell(w) &= \sum_{i=1}^n \left(y^{(i)} \ln \sigma(w \cdot x^{(i)}) + (1 - y^{(i)}) \ln(1 - \sigma(w \cdot x^{(i)})) \right) \\ w_{\text{LogR}} &\stackrel{\text{def.}}{=} \underset{w}{\operatorname{argmax}} \ell(w) = \underset{w}{\operatorname{argmin}} (-\ell(w)). \end{aligned} \quad (2)$$

$$P(Y = 1|X = x) = \sigma(z) \Leftrightarrow P(Y = 0|X = x) = 1 - \sigma(z), \text{ where}$$

$$\sigma(z) \stackrel{\text{def.}}{=} \frac{1}{1 + e^{-z}} = \frac{e^z}{1 + e^z},$$

$$z \stackrel{\text{not.}}{=} w_0 + \sum_{i=1}^d w_i x_i \stackrel{\text{not.}}{=} w \cdot x, \text{ with}$$

$$w \stackrel{\text{not.}}{=} (w_0, w_1, \dots, w_d) \in \mathbb{R}^{d+1}, \text{ assuming } x_0 = 1.$$

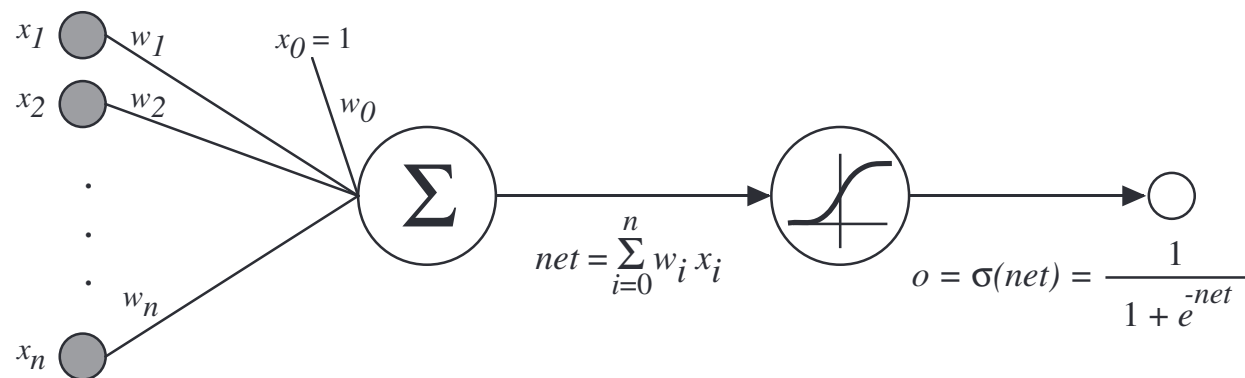


Note that $-\ell(w)$ is a cross-entropy.

For a more general result than (2), see Stanford, 2015 fall, Andrew Ng, HW 3, pr. 5.c.

See the analogy with the sigmoidal perceptron

CMU, 2011 fall, Eric Xing, HW1, pr. 3.3



The k -Nearest Neighbor Algorithm

Evelyn Fix, Joseph Hodges, 1951; Thomas Cover, Peter Hart, 1967

Training:

Store all training examples.

Classification:

Given a query/test instance x_q ,
first locate the k nearest training examples x_1, \dots, x_k ,
then estimate $\hat{f}(x_q)$:

- take a vote among its k nearest neighbors

$$\hat{f}(x_q) \leftarrow \operatorname{argmax}_{v \in V} \sum_{i=1}^k 1_{\{f(x_i)=v\}}$$

where $1_{\{.\}}$ is the well-known *indicator function*.

The Bottom-up Hierarchical Clustering Algorithm

Given: a set $X = \{x_1, \dots, x_n\}$ of objects
a function **sim**: $\mathcal{P}(X) \times \mathcal{P}(X) \rightarrow R$

for $i = 1, n$ **do**
 $c_i = \{x_i\}$ **end**
 $C = \{c_1, \dots, c_n\}$
 $j = n + 1$
while $|C| > 1$
 $(c_{n_1}, c_{n_2}) = \operatorname{argmax}_{(c_u, c_v) \in C \times C} \mathbf{sim}(c_u, c_v)$
 $c_j = c_{n_1} \cup c_{n_2}$
 $C = C \setminus \{c_{n_1}, c_{n_2}\} \cup \{c_j\}$
 $j = j + 1$

The k -Means Algorithm

S. P. Lloyd, 1957

Given: a set $X = \{x_1, \dots, x_n\} \subseteq \mathcal{R}^m$,
a distance measure d on \mathcal{R}^m ,
a function for computing the mean $\mu : \mathcal{P}(\mathcal{R}^m) \rightarrow \mathcal{R}^m$,

built k clusters so as to satisfy a certain (“stopping”) criterion (e.g., maximization of group-average similarity).

Procedure:

Select (arbitrarily) k initial centers f_1, \dots, f_k in \mathcal{R}^m ;

while the stopping criterion is not satisfied

for all clusters c_j **do** $c_j = \{x_i \mid \forall f_l \ d(x_i, f_j) \leq d(x_i, f_l)\}$ **end**

for all means f_j **do** $f_j \leftarrow \mu(c_j)$ **end**

K-Means algorithm revisited (I)

- Se inițializează în mod arbitrar centrozii $\mu_1, \mu_2, \dots, \mu_K$ și se ia $C = \{1, \dots, K\}$.
- Atâta timp cât valoarea criteriului J descreește în mod strict, repetă:

Pasul 1:

Calculează γ astfel:

$$\gamma_{ij} \leftarrow \begin{cases} 1, & \text{dacă } \|x_i - \mu_j\|^2 \leq \|x_i - \mu_{j'}\|^2, \forall j' \in C, \\ 0, & \text{în caz contrar.} \end{cases}$$

În caz de egalitate, alege în mod arbitrar cărui cluster (dintre cele eligibile) să-i aparțină x_i .

Pasul 2:

Recalculează μ_j folosind matricea γ actualizată:

Pentru fiecare $j \in C$, dacă $\sum_{i=1}^n \gamma_{ij} > 0$, asignează

$$\mu_j \leftarrow \frac{\sum_{i=1}^n \gamma_{ij} x_i}{\sum_{i=1}^n \gamma_{ij}}.$$

Altfel, menține neschimbat centroidul μ_j .

Algoritmul de clusterizare K -means poate fi văzut [și reformulat] ca un *algorithm de optimizare*, folosind **metoda descreșterii pe coordonate**.

Obiectivul este acela de a minimiza o funcție obiectiv care măsoară (indirect) coeziunea intra-clustere:

$$J(L, \mu) = \sum_{i=1}^n \|x_i - \mu_{l_i}\|^2,$$

Algoritmul K -means face *inițializarea* centroizilor clusterelor μ cu anumite valori, după care procedează astfel:

Pasul 1:

Păstrând μ fixat, găsește acea asignare L a instanțelor la clustere care minimizează funcția $J(L, \mu)$;

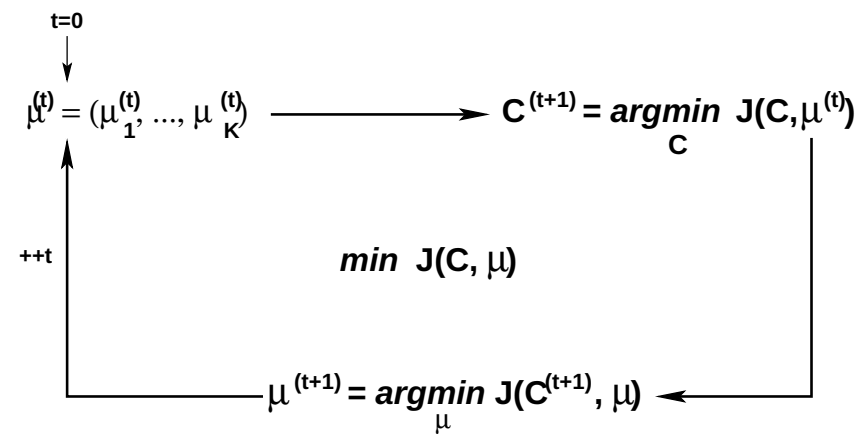
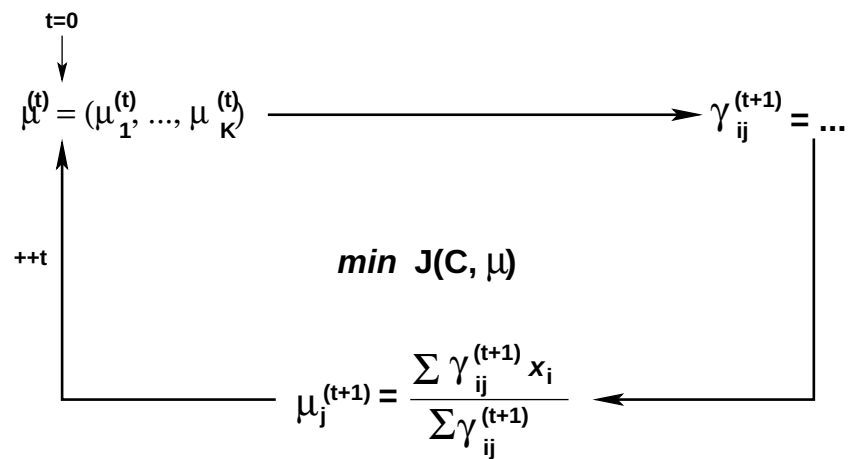
Pasul 2:

Păstrând asignarea L fixată, găsește acea valoare pentru μ pentru care se minimizează $J(L, \mu)$.

Criteriul de oprire: Dacă [aceasta nu este prima iterație și] niciuna dintre asignările din lista L nu s-a modificat în raport cu precedenta iterație, se trece la pasul următor (Terminare); altfel se repetă de la Pasul 1.

Terminare: Returnează L și μ .

K-Means algorithm revisited (II)



The General EM Problem

Approach

Given

- observed data $X = \{x_1, \dots, x_m\}$ independently generated using the parameterized distributions/hypotheses h_1, \dots, h_m
- unobserved data $Z = \{z_1, \dots, z_m\}$

determine

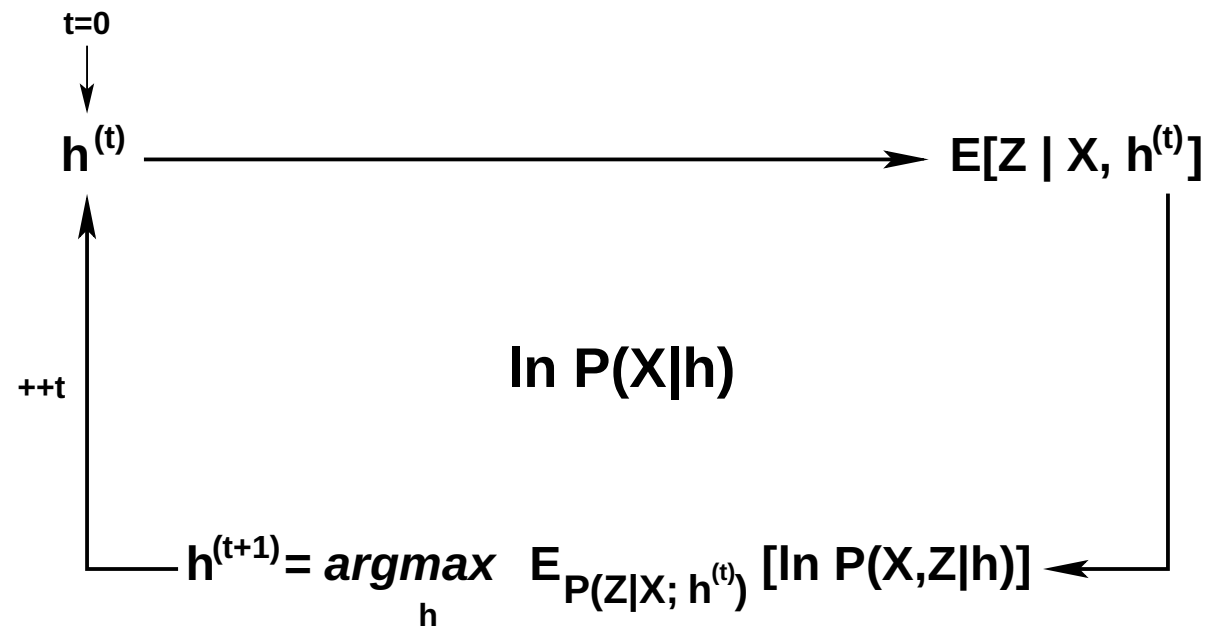
\hat{h} that (locally) maximizes $P(X|h)$.

Start with $h^{(0)}$, an arbitrarily/conveniently chosen value of h .

Repeatedly

1. Use the observed data X and the current hypothesis $h^{(t)}$ to **estimate [the probabilities associated to the values of] the unobserved** variables Z , and further on compute their expectations, $E[Z]$.
2. The expected values of the unobserved variables Z are used to **calculate an improved hypothesis** $h^{(t+1)}$, based on **maximizing the mean of a log-likelihood function**: $E[\ln P(Y|h)|X, h^{(t)}]$, where $Y = \{y_1, \dots, y_m\}$ is the complete (observed and unobserved) data, i.e. $y_i = (x_i, z_i)$, for $i = 1, \dots, m$.

The EM algorithmic Schema



ADMINISTRATIVA

Who is Liviu Ciortuz?

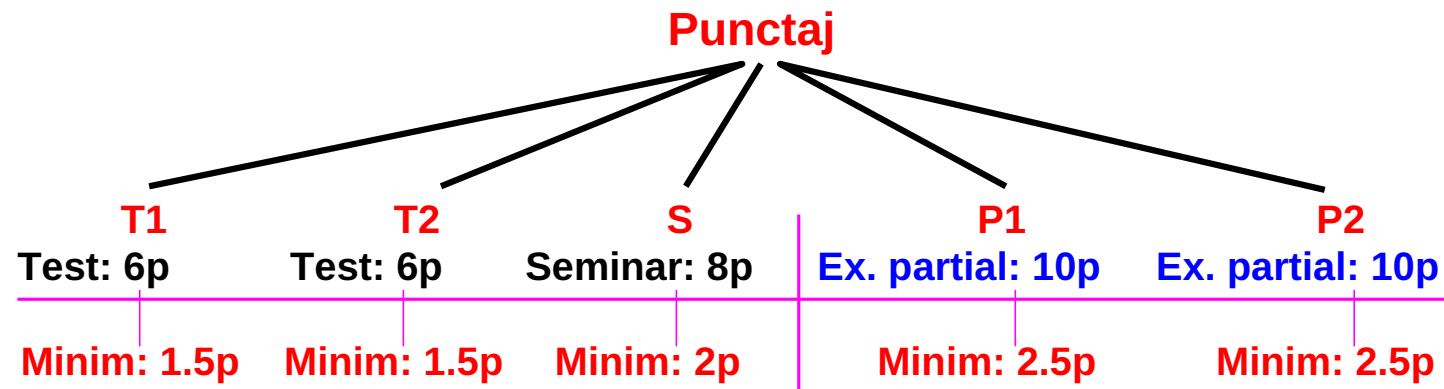
- Diploma (maths and CS) from UAIC, Iași, Romania, 1985
PhD in CS from Université de Lille, France, 1996
- programmer:
Bacău, Romania (1985-1987)
- full-time researcher:
Germany (DFKI, Saarbrücken, 1997-2001),
UK (Univ. of York and Univ. of Aberystwyth, 2001-2003),
France (INRIA, Rennes, 2012-2013)
- assistant, lecturer, and then associate professor:
Univ. of Iași, Romania (1990-1997, 2003-2012, 2013-today)

Teaching assistants for the ML undergraduate course 2023 (fall semester)

- Conf. dr. Anca Ignat (... Image processing)
<https://profs.info.uaic.ro/~ancai/ML/>
- Sebastian Ciobanu (PhD; Amazon)
<https://sites.google.com/view/seminarml>
- Andi Munteanu
(PhD student at UAIC, and research assistant at Univ. of Cambridge)
- Cristian Simionescu (PhD student; Nexus)
- Ramona Albert (PhD student; Amazon)
- Ștefan Panțiru (MSc; Mambu)
- Corina Dimitriu (MSc student)

Grading standards for the ML undergraduate course 2023

Obiectiv: Învățare pe tot parcursul semestrului!



Prezenta la curs: recomandata!

Prezenta la seminar: obligatorie!

Penalizare: 0.2p pentru fiecare absenta de la a doua incolo!

Nota = $(10 + T1 + T2 + S + P1 + P2) / 5$

Pentru promovare: Nota $\geq 4.5 \iff T1 + T2 + S + P1 + P2 \geq 12.5$

REGULI generale pentru cursul de Învățare automată de la licență

Regulile de organizare a cursului de Învățare Automată (engl., Machine Learning, ML), sem. I, sunt specificate în *fișa disciplinei*

<http://profs.info.uaic.ro/~ciortuz/fisa-disciplinei.pdf>

- Bibliografie minimală: vezi slide #8
- Planificarea materiei, pentru fiecare săptămână (curs + seminar):
<http://profs.info.uaic.ro/~ciortuz/what-you-should-know.pdf>
- Prezența la curs: recomandată!
- **Regula 0: Prezența la seminar: obligatorie!**

Pentru fiecare absență la seminar, începând de la a doua absență încolo, se aplică o penalizare/depunctare de 0.2 puncte. (Vezi formula de notare.)

Regulile se aplică inclusiv studenților reînmatriculați.

- Săptămânal — marțea, între orele 18–20, în sala C308 — se va ține un seminar suplimentar, destinat pentru acei studenți care sunt foarte interesați de acest domeniu și cărora le plac demonstrațiile matematice. (Vedeți secțiunile “Advanced issues” din documentul <http://profs.info.uaic.ro/~ciortuz/what-you-should-know.pdf>.)

REGULI generale pentru cursul de Învățare automată de la licență (cont.)

Regula 1: Pentru **seminarii**, nu se admit mutări ale studenților de la o grupă la alta, decât în cadrul grupelor care au același asistent / profesor responsabil de seminar.

Regula 2: Nu se fac **echivalări** de punctaje pentru studenții care nu au promovat cursul în anii precedenți.

Regula 3: Profesorul responsabil pentru acest curs, Liviu Ciortuz,
NU va răspunde la email-uri care pun întrebări pentru care răspunsul a fost deja dat

- fie în aceste slide-uri,
- fie pe **site-ul Piazza dedicat acestui curs:**
<https://piazza.com/info.uaic.ro/fall2023/ml2023f/home>,
- fie la curs.

Recomandare importantă (1) La fiecare curs și seminar, studenții vor avea **culegerea** de *Exerciții de învățare automată* (de L. Ciortuz et al) — vă recomandăm să imprimați capitolele *Clasificare bayesiană*, *Învățare bazată pe memorare*, *Arbori de decizie* și *Clusterizare* — și eventual **slide-urile** indicate în slide-ul următor.

Recomandare importantă (2) **Consultați săptămânal** documentul *what-you-should-know.pdf* din pagina de Resurse, de pe site-ul Piazza dedicat acestui curs.

REGULI generale pentru cursul de Învățare automată de la licență (cont.)

- **Slide-uri de imprimat** (în această ordine și, de preferat, COLOR):

<http://profs.info.uaic.ro/~ciortuz/SLIDES/foundations.pdf>

<https://profs.info.uaic.ro/~ciortuz/ML.ex-book/SLIDES/ML.ex-book.SLIDES.ProbStat.pdf>

<https://profs.info.uaic.ro/~ciortuz/ML.ex-book/SLIDES/ML.ex-book.SLIDES.DT.pdf>

<https://profs.info.uaic.ro/~ciortuz/ML.ex-book/SLIDES/ML.ex-book.SLIDES.Bayes.pdf>

<https://profs.info.uaic.ro/~ciortuz/ML.ex-book/SLIDES/ML.ex-book.SLIDES.IBL.pdf>

<https://profs.info.uaic.ro/~ciortuz/ML.ex-book/SLIDES/ML.ex-book.SLIDES.Cluster.pdf>

(Atenție: acest set de slide-uri poate fi actualizat pe parcursul semestrului!)

-
- **De imprimat (ALB-NEGRU):**

<http://profs.info.uaic.ro/~ciortuz/SLIDES/ml0.pdf>

<http://profs.info.uaic.ro/~ciortuz/SLIDES/ml3.pdf>

<http://profs.info.uaic.ro/~ciortuz/SLIDES/ml6.pdf>

<http://profs.info.uaic.ro/~ciortuz/SLIDES/ml8.pdf>

<http://profs.info.uaic.ro/~ciortuz/SLIDES/cluster.pdf>